

Arnab K. Paul

Postdoctoral Research Associate, Computing & Computational Sciences
Oak Ridge National Laboratory, U.S.A.

Oak Ridge, Tennessee
☎ +1 (540) 998-1480
✉ paula@ornl.gov
📄 [arnabkrpaul.github.io/](https://github.com/arnabkrpaul)

Research Interest

My interests lie in various domains of computer systems including **distributed systems**, **parallel file systems**, **big data analysis**, **edge computing**, and **high performance computing**. Recently, my interest has also piqued in the direction of **applying machine learning** to optimize system performance.

Education

- 2015–2020 **Ph.D., Computer Science, Virginia Tech.**
Dissertation: An application-attuned framework for optimizing HPC storage systems.
Advisor: Ali R. Butt. **GPA:** 4.0/4.0
- 2015–2018 **M.S., Computer Science and Applications, Virginia Tech.**
GPA: 3.85/4.0
- 2013–2015 **M.Tech., Computer Science and Engineering, National Institute of Technology, Rourkela.**
Thesis: Dynamic virtual machine placement in cloud computing.
Advisor: Bibhudatta Sahoo. **GPA:** 9.56/10.0 (Gold Medalist - First position in the department)
- 2009–2013 **B.Tech., Computer Science and Engineering, West Bengal University of Technology.**
GPA: 9.02/10.0 (First position in the department)
- 2009 **Class XII - I.S.C. Examination, Council for The Indian School Certificate Examination.**
Hill Top School, Jamshedpur
Marks: 96.25%
- 2007 **Class X - I.C.S.E, Council for The Indian School Certificate Examination.**
Hill Top School, Jamshedpur
Marks: 95.8%

Research Experience

- 09/20–present **Oak Ridge National Laboratory - AI Methods at Scale Group, Postdoctoral Research Associate.**
 - Analyze the I/O patterns of emerging data science applications, and identify performance bottlenecks.
 - Develop an efficient cache for exascale deep learning applications.
- 08/15–08/20 **Virginia Tech - Distributed Systems and Storage Laboratory, Ph.D. Student in Dept. of CS.**
 - Conducted an empirical study on the use of containers in HPC platforms.
 - Developed an approach for estimating the performance of edge-based clustering applications.
 - Built an I/O framework for load balancing storage servers in HPC parallel file systems, like Lustre.
 - Developed a model to optimize data partitioning for in-memory data analytics platforms, like Spark.
- 06/19–08/19 **Cray Inc., Graduate Research Intern.**
Mentors: Cory Spitz, Nathan Rutman (Cray Inc.), and Scott White (Los Alamos National Laboratory)
 - Built a scalable re-indexer for BRINDEXER - a metadata indexing tool used in Cray.
- 05/18–08/18 **Lawrence Livermore National Laboratory, Graduate Student Summer Intern(Computation Scholar).**
Mentor: Kathryn Mohror
 - Analyzed the characteristics of metadata and I/O for jobs running on two supercomputers at LLNL.
 - Built an I/O scheduler for scheduling I/O for jobs to reduce I/O contention.
- 05/17–08/17 **Argonne National Laboratory, Graduate Student Summer Intern(Research Aide).**
Mentor: Ian Foster
 - Created FSMonitor - a tool for scalable file system event monitoring for arbitrary file systems.
- 01/14–05/15 **NIT Rourkela - Information Security and Data Communication Laboratory, M.Tech. Student.**
 - Proposed an approach for dynamic virtual machine placement in the cloud using game theory.
 - Applied and analyzed greedy algorithms on virtual machine distribution across data centers.

Awards and Honors

- 2020 Nominated for the YESC award for most innovative student paper at IEEE Services 2020, Beijing, China
- 2019-2020 BitShares Fellowship, Department of Computer Science, Virginia Tech
- 2019 Travel Grant Recipient, IEEE Cluster, Albuquerque, NM, USA
- 2019 Student Volunteer, SC, Denver, CO, USA
- 2018, '19, '20 Member of the Dean's Graduate Team & Ambassador to the College of Engineering, Virginia Tech
- 2018 Student Volunteer, SCiNet @ SC, Dallas, TX, USA
- 2017-2018 President, Bengali Students' Association, Virginia Tech
- 2017-2018 BitShares Fellowship, Department of Computer Science, Virginia Tech
- 2017 Travel Grant Recipient, IEEE BigData, Boston, MA, USA
- 2016 Travel Grant Recipient, IEEE Cluster, Taipei, Taiwan
- 2016 Student Volunteer, SC, Sal Lake City, Utah, USA
- 2015 Gold Medalist, Department of Computer Science, National Institute of Technology - Rourkela
- 2013 Recognition for Rank 1, Department of Computer Science, West Bengal University of Technology
- 2009 Recognition for holding 1st rank from Kindergarten to Std. XII (all 15 years), Hill Top School

Teaching Experience

- Fall 2019 **Instructor, Virginia Tech** ↗ courses.cs.vt.edu/cs2505/fall2019/.
CS2505: Introduction to Computer Organization - I: Prepared and gave lectures to two sections (~150 students), prepared assignments and examinations, awarded final grades, mentored graduate teaching assistants.
- 2015–2019 **Graduate Teaching Assistant, Department of Computer Science, Virginia Tech.**
- Spring 2019 CS 3214: Operating Systems *Recitation sessions, grading, guest lectures, office hours*
- Fall 2018 CS 5584: Network Security *Project ideas with 15 groups, grading, office hours*
- Spring 2018 CS 3114: Data Structures and Algorithms *Grading, office hours*
- Fall 2017 CS 2506: Introduction to Computer Organization - II *Grading, office hours*
- Spring 2017, CS 2114: Software Design and Data Structures *Lab sessions for 60 students, practice sessions,*
Fall 2016 *designing and grading assignments, office hours*
- Spring 2016, CS 1054: Introduction to Programming in Java *Lab sessions for 60 students, grading, office hours*
Fall 2015
- 2014–2015 **Graduate Teaching Assistant, Department of Computer Science, NIT Rourkela.**
- Autumn 2014, CS 171: Computing Lab *Prepared and gave lectures, held lab sessions for 220 students,*
Spring 2015 *preparing and grading assignments*
- Spring 2015 CS 670: Data Mining Lab *Lab sessions for 30 students, grading*

Publications

Book Chapters

- CRC Press '20 **Arnab K. Paul**. Edge or Cloud: What to Choose?. In Cloud Network Management: An IoT based Framework, CRC Press, Taylor & Francis Group, pages 14, 2020.
- IGI Global '17 **Arnab Kumar Paul**, and Bibhudatta Sahoo. Dynamic virtual machine placement in cloud computing. In Resource Management and Efficiency in Cloud Computing Environments, pp. 136-167, IGI Global, 2017. ↗ doi.org/10.4018/978-1-5225-1721-4.ch006

Journal Publications

- TPDS '20 Nannan Zhao, Vasily Tarasov, Hadeel Albahar, Ali Anwar, Lukas Rupperecht, Dimitrios Skourtis, **Arnab K. Paul**, Keren Chen, and Ali R. Butt. Large-Scale Analysis of Docker Images and Performance Implications for Container Storage Systems. Accepted In IEEE Transactions on Parallel and Distributed Systems (TPDS), pages 12, 2020.
- [Core Rank: A*]

Conference Publications

- HiPC '20 [Core Rank: National - India] **Arnab K. Paul**, Olaf Faaland, Adam Moody, Elsa Gonsiorowski, Kathryn Mohror, and Ali R. Butt. Understanding HPC Application I/O Behavior Using System Level Statistics. Accepted In Proceedings of the 27th IEEE International Conference on High Performance Computing, Data, and Analytics, Pune, India, pages 10, December 2020.
- SMDS '20 Breno Dantas Cruz, **Arnab K. Paul**, Zheng Song, and Eli Tilevich. STARGAZER: A Deep Learning Approach for Estimating the Performance of Edge-Based Clustering Applications. In Proceedings of the IEEE International Conference on Smart Data Services, Beijing, China, pages 9, October 2020. (Nominated for the YESC award for the most innovative student paper at IEEE Services 2020.)
- Cloud '20 [Core Rank: B] Subil Abraham¹, **Arnab K. Paul**¹, Redwan Ibne Seraj Khan, and Ali R. Butt. On the Use of Containers in High Performance Computing Environments. In Proceedings of the IEEE International Conference on Cloud Computing, Beijing, China, pages 9, October 2020. (AR: 17%).
¹ Both authors contributed equally.
- CCGrid '20 [Core Rank: A] **Arnab K. Paul**, Brian Wang, Nathan Rutman, Cory Spitz, and Ali R. Butt. Efficient Metadata Indexing for HPC Storage Systems. In Proceedings of the 20th IEEE/ACM International Symposium on Cluster, Cloud and Internet Computing, Australia, pages 10, May 2020. (AR: 23%). doi.org/10.1109/CCGrid49817.2020.00-77
- Cluster '19 [Core Rank: A] **Arnab K. Paul**, Ryan Chard, Kyle Chard, Steven Tuecke, Ali R. Butt, and Ian Foster. FSMonitor: Scalable File System Monitoring for Arbitrary Storage Systems. In Proceedings of the IEEE International Conference on Cluster Computing, Albuquerque, NM, pages 11, September 2019. (AR: 22%). doi.org/10.1109/CLUSTER.2019.8891045
- IPDPS '19 [Core Rank: A] Bharti Wadhwa, **Arnab K. Paul**, Sarah Neuwirth, Feiyi Wang, Sarp Oral, Ali R. Butt, Jon Bernard, and Kirk W. Cameron. iez: Resource Contention Aware Load Balancing for Large-Scale Parallel File Systems. In Proceedings of the IEEE International Parallel and Distributed Processing Symposium, Rio de Janeiro, Brazil, pages 11, May 2019. (AR: 25%). doi.org/10.1109/IPDPS.2019.00070
- BigData '17 **Arnab K. Paul**, Arpit Goyal, Feiyi Wang, Sarp Oral, Ali R. Butt, Michael J. Brim, and Sangeetha B. Srinivasa. I/O Load Balancing for Big Data HPC Applications. In Proceedings of the IEEE International Conference on Big Data, Boston, MA, pages 10, December 2017. (AR: 18%). doi.org/10.1109/BigData.2017.8257931
- Cluster '16 [Core Rank: A] **Arnab Kumar Paul**, Wenjie Zhuang, Luna Xu, Min Li, Mustafa Rafique, and Ali R. Butt. CHOPPER: Optimizing Data Partitioning for In-Memory Data Analytics Frameworks. In Proceedings of the IEEE International Conference on Cluster Computing, Taiwan, pages 10, September 2016. (AR: 24%). doi.org/10.1109/CLUSTER.2016.41
- INDICON '14 **Arnab Kumar Paul**, Sourav Kanti Addya, Bibhudatta Sahoo, and Ashok Kumar Turuk. Application of greedy algorithms to virtual machine distribution across data centers. In Proceedings of 2014 Annual IEEE India Conference, pp. 1-6. IEEE, 2014. doi.org/10.1109/INDICON.2014.7030633
- ICACCCT '14 Arjun Datta, and **Arnab Kumar Paul**. Online compiler as a cloud service. In Proceedings of 2014 IEEE International Conference on Advanced Communications, Control and Computing Technologies, pp. 1783-1786. IEEE, 2014. doi.org/10.1109/ICACCCT.2014.7019416

Workshop Publications

- PDSW '19 @ SC '19 [Core Rank: A] **Arnab K. Paul**, Olaf Faaland, Adam Moody, Elsa Gonsiorowski, Kathryn Mohror, and Ali R. Butt. Improving I/O Performance of HPC Application Using Intra-Job Scheduling. Work-In-Progress in Proceedings of the 4th Joint International Workshop on Parallel Data Storage & Data Intensive Scalable Computing Systems ((PDSW-DISC'17)) in conjunction with SC'19, Denver, CO, pages 1, November 2019. www.pdsw.org/pdsw19/wips/ArnabPaul-pdswWIP.pdf
- ComNet-IoT @ ICDCN '19 [National - India] Hyogi Sim, **Arnab K. Paul**, Eli Tilevich, Ali R. Butt, and Muhammad Shahzad. CSLIM: Automated Extraction of IoT Functionalities from Legacy C Codebases. In Proceedings of the 8th International Workshop on Computing and Networking for IoT and Beyond in conjunction with ICDCN '19, Bangalore, India, pages 6, January 2019. doi.org/10.1145/3288599.3296013
- PDSW '17 @ SC '17 [Core Rank: A] **Arnab K. Paul**, Ryan Chard, Kyle Chard, Steven Tuecke, Ali R. Butt, and Ian Foster. Toward Scalable Monitoring on Large-Scale Storage for Software Defined Cyberinfrastructure. In Proceedings of the 2nd Joint International Workshop on Parallel Data Storage & Data Intensive Scalable Computing Systems (PDSW-DISC'17) in conjunction with SC'17, Denver, Colorado, pages 6, November 2017. doi.org/10.1145/3149393.3149402

WHPC '16 @ SC '16 [Core Rank: A] Sangeetha B. Srinivasa, **Arnab K. Paul**, Arpit Goyal, Feiyi Wang, Sarp Oral, and Ali R. Butt. I/O Load Balancing for Lustre Distributed File System. In Women in HPC in conjunction with SC'16, Salt Lake City, Utah, November 2016.

Posters

- SC '19 [Core Rank: A] **Arnab K. Paul**, Olaf Faaland, Adam Moody, Elsa Gonsiorowski, Kathryn Mohror, and Ali R. Butt. Understanding HPC Application I/O Behavior Using System Level Statistics. In Proceedings of The International Conference for High Performance Computing, Networking, Storage, and Analysis (SC) 2019, Denver, CO, pages 3, November 2019. sc19.supercomputing.org/proceedings/tech_poster/tech_poster_pages/rpost157.html
- Cray '19 **Arnab K. Paul**, Nathan Rutman, Cory Spitz, Brian Wang, Peter Bojanic, and Ali R. Butt. Analyzing the performance of file system indexing tools. In Cray Inc. Summer Student Poster Session, Minneapolis, MN, August 2019.
- LLNL '18 **Arnab K. Paul**, Olaf Faaland, Adam Moody, Elsa Gonsiorowski, Kathryn Mohror, and Ali R. Butt. Analysis and predictive modeling of HPC I/O workloads. In LLNL Computation Summer Student Poster Session, Livermore, CA, August 2018.

Theses

- Ph.D. '20 **Arnab Kumar Paul**. An application-attuned framework for optimizing HPC storage systems. Ph.D. dissertation, Department of Computer Science and Applications, Virginia Tech, U.S.A., 2020. hdl.handle.net/10919/99793
- M.Tech. '15 **Arnab Kumar Paul**. Dynamic virtual machine placement in cloud computing. Master's thesis, Department of Computer Science and Engineering, National Institute of Technology, Rourkela, 2015. ethesis.nitrkl.ac.in/6811/

Talks and Presentations

- 2020 **Oak Ridge National Laboratory**, *A Framework for Whole Stack Optimization of Distributed Storage Systems*, Job talk.
- 2020 **Lawrence Berkeley National Laboratory**, *A Framework for Whole Stack Optimization of Distributed Storage Systems*, Job talk.
- 2019 **SC**, *Understanding HPC Application I/O Behavior Using System Level Statistics*, Poster presentation.
- 2019 **SC**, *Improving I/O Performance of HPC Application Using Intra-Job Scheduling*, WIP presentation.
- 2019 **Cluster**, *Scalable File System Monitoring for Arbitrary Storage Systems*, Paper presentation.
- 2019 **ICDCN**, *Automated extraction of IoT functionalities from legacy C codebases*, Paper presentation.
- 2017 **BigData**, *I/O Load Balancing for Big Data HPC Applications*, Paper presentation.
- 2017 **PDSW @ SC**, *Toward Scalable Monitoring on Large-Scale Storage for Software Defined Cyberinfrastructure*, Paper presentation.
- 2016 **Cluster**, *Optimizing Data Partitioning for In-Memory Data Analytics Frameworks*, Paper presentation.

Mentoring Experience

- Graduate Students Redwan Ibne Seraj Khan (Ph.D., Virginia Tech, 2019 -)
Debasmita Biswas (MS, Virginia Tech, 2020 -)
Subil Abraham (MS, Virginia Tech, 2019 - 2020)
Arpit Goyal (MS, Virginia Tech, 2016 - 2017)
- Undergraduate Students Subrat Dhal (B.Tech., NIT Rourkela, 2014 - 2015)
Harshit Verma (B.Tech., NIT Rourkela, 2014 - 2015)

Professional Service

- TPC Member ICDCS '20
- Reviewer IEEE Transactions on Parallel and Distributed Systems (TPDS) '19 '20, Neural Processing Letters (NEPL) '20, Cluster Computing Journal '19 '20, IJGHP '18 '19 '20, ASTESJ '18, AUTOSOFT Journal '18, MGS Journal '17

External IEEE TSC Journal '18, BigData '17 '18, Cluster '17 '18 '20, ECOOP '20, HPDC '17 '18 '20, IC2E '17,
Reviewer ICCD '19, ICDCS '17 '18 '19, ICS '17 '18, IPDPS '18 '19 '20
Facilitator SC '20, HPDC '19

Skills

General C, C++, Python, Java, UNIX, git, svn, latex, gnuplot.
Analytics Apache Spark, pandas, matplotlib, bigdata analysis, applied machine learning, federated learning.
File Systems Lustre file system, Ceph object store, HDFS, IBM Spectrum Scale (GPFS).
Distributed Containers, cloud computing, key-value stores, edge computing, IoT, map-reduce.
Computing

Memberships

2019 – present Association for India's Development, Blacksburg Chapter.
2018 – 2020 Graduate Students' Council, Department of Computer Science, Virginia Tech.
2016 – present Institute of Electrical and Electronics Engineers (IEEE), Student member.

References

1. Dr. Ali R. Butt <butta@vt.edu>, *Professor, Virginia Tech.*
2. Dr. Ian Foster <foster@anl.gov>, *Senior Scientist - Argonne National Laboratory, Professor - University of Chicago.*
3. Dr. Eli Tilevich <tilevich@cs.vt.edu>, *Associate Professor, Virginia Tech.*
4. Dr. Bibhudatta Sahoo <bdsahu@nitrkl.ac.in>, *Associate Professor, National Institute of Technology, Rourkela.*